

Postgres Local First

App code works directly with client-side embedded database which auto syncs with backend database

Reads & writes go to local DB first

Benefits

App devs	End users
Simple state mgmt	Instant reactive UX
Simple backend	Offline first
↓ backend compute	Real time multi-user collaboration

Primarily use Postgres & Solite

Solite is hard to beat on client-side

→ Battle tested (Trillion Solite DB in active use)

→ Perf & aggregations

Objectives

→ Treat PG DB as sacrosanct

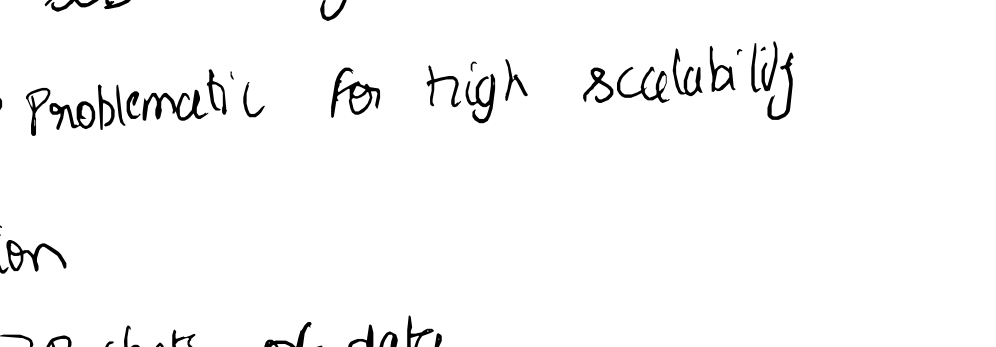
→ Min changes to schema, config

→ Don't bypass user logic

→ Dynamic partial replication

Architecture

Separate Read/Download path from write/upload path



CRDT is key

Implementing dynamic partial Replication

→ Make arbitrary queries & keep results sets in sync

→ Problematic for high scalability

Solution

→ Buckets of data

→ Shared between clients

→ Changes to buckets are synced

Based on sync rules, replicate & pre-process data from PG

First take snapshot & then update incrementally (using logical replication)

→ Checkpoints keep track of LSNs

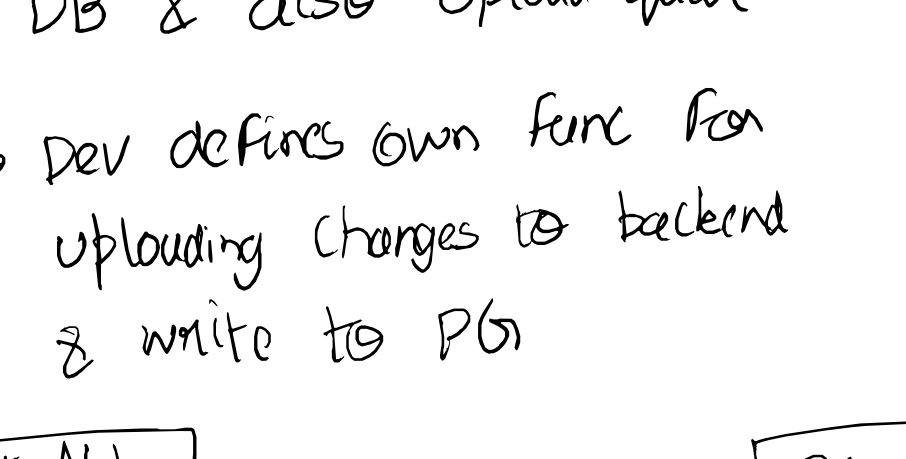
→ Store recent history of ops for each row

→ Rep as JSON

→ Ops indexed by 'operation ID' which is strictly ↑

→ Allows query of ordered ops efficiently

3. Auth users using JWTs



4. Streaming sync of bucket data from service to clients

5. On client, data is persisted to Solite

→ Type mapping from PG → Solite

→ Replicated data stored in schemaless format

→ client side schema

→ Schema is applied as Solite views on top of schemaless data

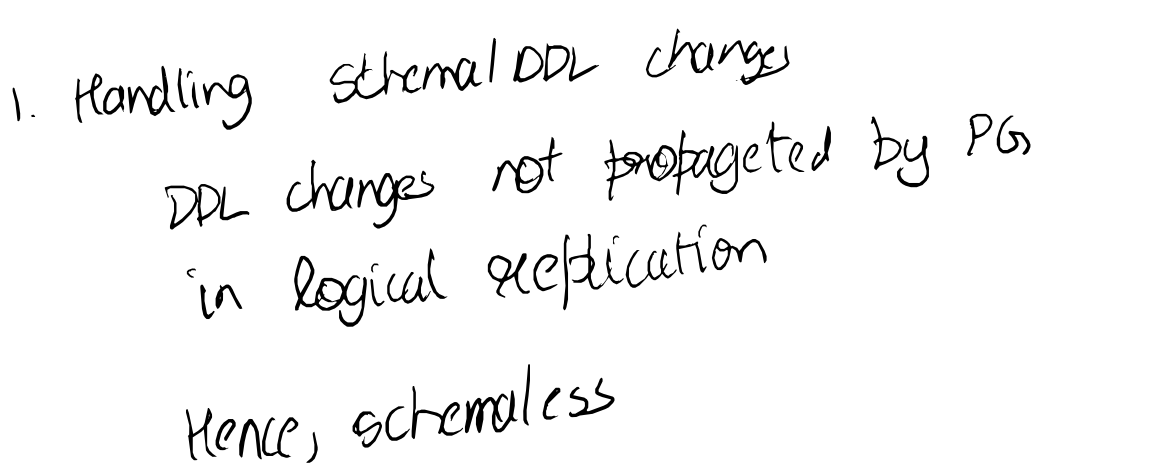
→ Live reactive query hooks:

Update UI when data changes (React hooks)

6. Syncing local changes

→ Changes written to local Solite DB & also upload queue

→ Dev defines own func for uploading changes to backend & write to PG



Server authentication: NO CRDT's needed

CRDT → Special datatype where changes can be merged in any order & resolution is always deterministic

However, NO CRDT's used in this version:

Alternative: Always merge changes in same order

Guaranteeing consistency

→ Track write ckpts on server

→ Ckpts have op-id & Postgres LSN

→ Local client side writes to Solite database

→ Applied on top of last read ckpt from server

→ Added to upload queue

→ Client retrieves latest ckpt from server AFTER it has finished uploading writes

→ Client update local state to match server state.

Guaranteeing data integrity

→ Checksum computation on bucket.

When mismatch, download from bucket

Problems & solution

1. Handling schema/DDI changes

DDL changes not propagated by PG in logical replication

Hence, schemaless

DDL → CREATE, DROP, RENAME, REPLICATION

2. LSN's that overlap btw transactions

→ LSN's are not guaranteed to be monotonically increasing

→ Sharded DB support

→ Reprocess sync data

3. Accessing TOASTed values

4. Handling diff row id permutations

Implemented support for diff permutations

